

(a) Image-pretrained foundation



Image pretraining

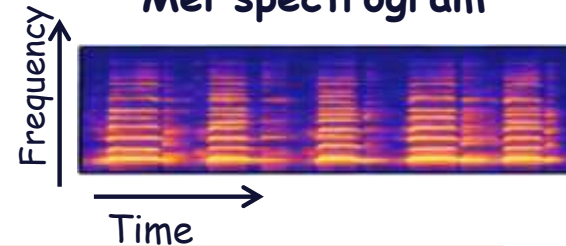


FLUX.2-klein-base-4B

Transfer weights

(b) Reuse the image pipeline

Mel spectrogram



Pseudo-RGB

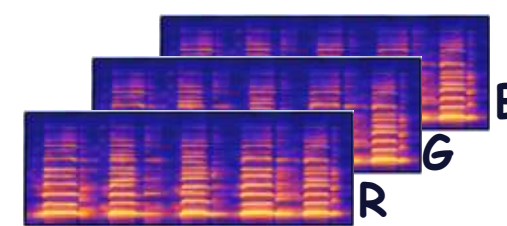
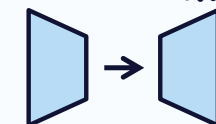


Image DiT



Fine-tuned on speech

Image VAE



Frozen

Qwen3 text encoder

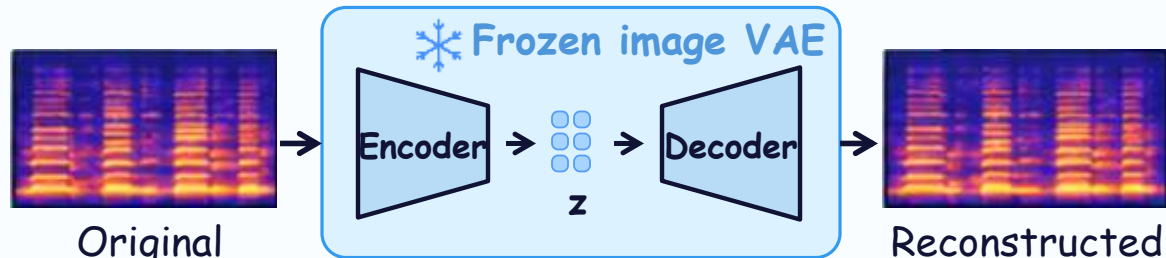


Frozen

Original architectures retained

(c) Three contributions

1 Speech-compatible image latents



High-fidelity reconstruction

2 Transferable image priors



Faster convergence

Image-pretrained initialization



Better results

WER at 100k updates

4.2%

Random initialization



2.7%

Image-pretrained

3 Voice and style control

Zero-shot TTS



New text



Reference speech



Same voice, new words

Instruction-guided TTS

Emotion

Rate

Accent



Controlled speech

Speak warmly and slowly